

The Consumer-Relative Flip Across Domains: Preserve What the Reader Reads

Anonymous authors

Paper under double-blind review

Abstract

Signals are almost universally compressed to minimise reconstruction error, on the assumption that a more faithfully reconstructed signal better serves whatever reads it next. For a consumer that reads only part of what a signal carries, this assumption is directionally wrong: downstream error is governed not by the total reconstruction error $\text{tr}(\Sigma_\delta)$ but by its projection onto the consumer’s *read operator*, $\text{tr}(P_C \Sigma_\delta)$. A code that spends a fixed budget minimising reconstruction spends it largely on directions the consumer ignores; a code that instead protects the read directions reconstructs the signal *worse* yet serves the consumer *better* — a dissociation we call the **flip**. We test the flip under a single pre-registered protocol across twelve domains spanning three distinct physics: direction finding on automotive radar, microphone, and seismic arrays; the attention keys of a language model; and learned classifiers of moral judgment, music genre, and sperm-whale dialect. The flip reproduces wherever the consumer’s read operator is **identifiable** — given by the geometry, or recovered blind from the consumer by probing what its output is sensitive to — and fails, informatively, exactly where the read operator is *coupled* to the signal’s own energy or the consumer does not function. We derive the failure modes as the four degeneracies of P_C and give the central one an ex-ante *sign* predictor: a read-energy alignment coefficient $\kappa = \text{tr}(\bar{P}_C \Sigma_x) / \sum_{i \leq r} \lambda_i(\Sigma_x)$ that near 1 forbids the flip before any code is built (the gradient-coupling null). We tested, and report as an honest negative, the stronger claim that κ reads off the flip *magnitude*: it does not — the magnitude is confounded by read-subspace energy and the quantizer regime, so κ is a coupling-null predictor, not a continuous dial. Which code best exposes the flip is set not by one statistic but by two axes of the representation’s geometry — concentration and correlation — a diagnostic that also selects the code at deployment and whose single-statistic shortcut we pre-register and refute. Reconstruction fidelity is the special case of a consumer that reads everything; observation, in general, is selective — and compression should be too.

Tiers, worn in public. Seven sealed confirmations (across synthetic, LLM-attention, acoustic, seismic, retrieval, and bioacoustic consumers) carry the claim; one partial (radar, data-limited), one sealed null (gradient compression, the coupling boundary), and two (A2)-probe verdicts (music, attention keys) are labelled as such throughout. Every row resolves to a sealed pre-registration and a result artifact.

1 Introduction

Almost every system that compresses a signal is judged by how faithfully it can reconstruct it — mean-squared error for images, cosine similarity for embeddings, PSNR for audio. The assumption underneath is so ingrained it is rarely stated: that a code which reconstructs the signal better is a code that serves whatever reads the signal next better. For a decoder that consumes only part of what the signal carries, this assumption is not merely loose — it is *directionally wrong*.

The quantity that governs a consumer’s downstream error is not the total reconstruction error $\text{tr}(\Sigma_\delta)$ but its projection onto the consumer’s **read operator**, $\text{tr}(P_C \Sigma_\delta)$ — the error in the directions the consumer

actually reads. When a fixed bit budget is spent to minimise reconstruction, it is spent largely on directions the consumer ignores; a code that instead protects the read directions reconstructs the raw signal *worse* yet serves the consumer *better*. We call this dissociation the **flip**: at matched bits, read-preserving coding beats reconstruction-optimal coding on the downstream task, precisely while reconstruction-optimal coding wins on reconstruction. The two fidelities move in opposite directions.

The flip was first isolated in two narrow settings — the direction of arrival read by a sensor array, and the keys read by softmax attention in a language model — where it could be dismissed as an artifact of the geometry. Here we ask whether it is instead a **property of observation**. We test the same matched-bits dissociation, with the same discipline (each configuration frozen on a calibration split and scored on held-out data, every outcome pre-registered and reported regardless of sign), across twelve domains spanning three distinct physics — direction finding on automotive radar, hearing-aid microphone arrays, and a seismic array reading earthquake back-azimuth; the attention keys of a transformer; a fine-tuned moral-judgment classifier; music genre and, strikingly, the dialect of sperm whales in their click-rhythm codas.

The dissociation reproduces across all three physics and every consumer type wherever a single condition holds: the **read operator is identifiable** — given by the geometry, or recovered blind from the consumer by probing what its output is sensitive to. Where it is not identifiable, or where it is *coupled* to the signal’s own energy (as it is for gradient compression under a curvature metric, whose high-curvature and high-energy directions coincide), the flip provably cannot appear — and it does not. These failures are not noise; they are the theory’s scope conditions. We show they are exactly the four ways the operator $P_{\mathcal{C}}$ can degenerate, and that the central one — coupling — is read off before any code is built by a read-energy alignment coefficient κ that near 1 forbids the flip (though, as we report, it does not predict the flip’s magnitude). Which encoding best exposes the flip, in turn, lives on two axes — how *concentrated* a representation’s directions are and how *correlated* its channels are — and no single summary statistic (we pre-register one and refute it) can stand in for the full two-axis diagnostic.

Contributions.

1. The flip demonstrated as **domain-general** — seven sealed confirmations spanning synthetic, LLM-attention, acoustic, seismic, retrieval, and bioacoustic consumers, three distinct physics — under sealed pre-registration with outcomes reported regardless of sign.
2. A **blind read-operator probe** that recovers what a consumer reads without its ground truth, turning a failed retrieval test into a confirmation on the identical held-out set, and confirmed a second time on a fully virgin split — the identifiability boundary made constructive.
3. The failure modes derived as the **four degeneracies of $P_{\mathcal{C}}$** , and the coupling degeneracy given an ex-ante **sign predictor** κ that forbids the flip when the read subspace aligns with the signal’s energy — together with the honest negative that κ does *not* predict the flip magnitude, which we tested and report rather than claim.
4. A **two-axis regime map** (concentration $\times$ correlation) that predicts the winning code per domain, with a validated diagnostic and a refuted single-statistic shortcut.

2 Theory

2.1 The controlling quantity and the flip

We take as given the second-order result established in a companion foundational work (Anonymous, 2026a): for a consumer $\mathcal{C}$ with output metric G acting on a representation through a Jacobian J , the leading-order downstream distortion induced by a coding error δ is

$$d_{\mathcal{O}} = \mathbb{E}_x \text{tr}(P_{\mathcal{C}}(x) \Sigma_{\delta}(x)), \quad P_{\mathcal{C}} = \mathbb{E}[J^{\top} G J], \quad (1)$$

reducing under stated conditions to $\text{tr}(P_{\mathcal{C}} \Sigma_{\delta})$, with Σ_{δ} the error second moment. $P_{\mathcal{C}}$ is symmetric positive semidefinite; its range is the *read subspace* (the directions the consumer distinguishes) and its kernel is the

quotient the consumer cannot tell apart. Reconstruction fidelity is the special case $P_C = I$: a consumer that reads everything.

Fix a total bit budget and consider three codes on the representation: a *read-preserving* code (R) that allocates bits to the range of P_C ; a *reconstruction-optimal* code (O) that minimises $\text{tr}(\Sigma_\delta)$ (reverse water-filling (Cover & Thomas, 2006) in the signal eigenbasis); and an *anti* code (A) that starves the read directions. The **flip** is the joint event

$$\underbrace{d_C(\text{R}) \leq d_C(\text{O})}_{\text{downstream: R wins}} \quad \text{and} \quad \underbrace{\text{tr}(\Sigma_\delta^{\text{O}}) \leq \text{tr}(\Sigma_\delta^{\text{R}})}_{\text{reconstruction: O wins}}, \quad (2)$$

with the anti check $d_C(\text{A}) \geq \max(d_C(\text{R}), d_C(\text{O}))$. The two inequalities in (2) point in opposite directions; that is the whole content. A code can win reconstruction and lose the task.

2.2 The four degeneracies of P_C

The flip in (2) is not automatic. It requires the reconstruction-optimal allocation to *differ* from the read-preserving one — i.e. the read subspace must be misaligned with the signal’s high-energy directions — and it requires the read operator to be correctly known. There are exactly four ways this can fail, and each is a degeneracy of P_C (or of its estimate $\hat{P}_C$):

- (D1) **Coupling** — range P_C **aligned with source energy**. If the read subspace coincides with the top-energy eigenspace of Σ_x , then O already allocates bits to the read directions and R has nothing to add. The flip vanishes although P_C is exact. This is a genuine wall, not an artifact; it is the empirical content of the gradient-compression null (§4). We make it a continuous coordinate in §2.3.
- (D2) **Precondition** — $P_C \approx 0$. If the consumer does not function, there is no read subspace to protect and (1) is degenerate. Observed for moral judgment on frozen embeddings (chance accuracy): the flip is ill-posed until the consumer is trained.
- (D3) **Identifiability** — **wrong $\hat{G}$ (hence wrong $\hat{P}_C$)**. If the read operator is *estimated* rather than recovered from the consumer, the estimate can protect the wrong subspace and O wins on held-out data. This is the only fixable failure: recovering P_C from the consumer’s own output sensitivity (§2.4) restores the flip.
- (D4) **Mechanism-absence** — **the claimed nuisance is not in $\ker P_C$** . If a direction hypothesised to be invariant is not actually in the consumer’s kernel, protecting it buys nothing and the effect does not exist. A genuine refutation, not a rehabilitatable miss.

A demonstration shows only the wins; a theory states where it fails. The four degeneracies are what let us say, before any experiment, which domains admit the flip.

2.3 The alignment law: coupling as a continuous dial

Degeneracy (D1) is the informative one, because it is not binary. Let $\bar{P}_C$ be the orthogonal projector onto the rank- r read subspace range P_C , and let $\lambda_1 \geq \dots \geq \lambda_d$ be the eigenvalues of the signal second moment Σ_x .

Definition 1 (read-energy alignment).

$$\kappa = \frac{\text{tr}(\bar{P}_C \Sigma_x)}{\sum_{i=1}^r \lambda_i(\Sigma_x)} \in \left(\frac{\sum_{i>d-r} \lambda_i}{\sum_{i \leq r} \lambda_i}, 1 \right]. \quad (3)$$

κ is the fraction of the maximal r -dimensional signal energy that the read subspace actually captures. $\kappa = 1$ iff the read subspace is the top- r energy eigenspace of Σ_x ; κ is small when the read subspace lies in the signal’s low-energy directions.

Proposition 1 (coupling null). At $\kappa = 1$ (full coupling, D1) the read subspace is the top-energy subspace, reverse water-filling already spends its budget on the read directions, and the flip vanishes: $\Delta = d_C(\text{O}) - d_C(\text{R}) \rightarrow 0$. Away from coupling, Δ can be positive.

Remark (what κ does and does not predict — honest negative, with a sharpening). The natural strengthening — that Δ is *monotone* decreasing in κ , so that κ reads off the flip magnitude as a single ex-ante number — is **not supported**, and pursuing the salvage yields a precise reason rather than another failure. We built the faithful coupling model (concentrated flat-block signal energy, read dimension equal to the energy-block dimension, and the scale-free relative flip $\Delta_{\text{rel}} = 1 - d_{\mathcal{C}}(\text{R})/d_{\mathcal{C}}(\text{O})$) and swept κ from 1 to 0. Two facts emerge. (a) *The coupling null is exact under an identifiable condition*: $\Delta_{\text{rel}} \rightarrow 0$ at $\kappa = 1$ iff the signal support equals the read support (the background energy vanishes) — which is precisely the gradient-compression condition, where the Hessian’s support coincides with the $X^{\top}X$ energy. (b) *In that same sharp-null regime the sweep degenerates*: once the background energy is small enough to make the null sharp, dialing κ downward moves the read subspace into the vanishing-energy background, so $\Delta_{\text{rel}} \approx 0$ for all $\kappa > 0$ and jumps only at $\kappa = 0$ — the flip lands on negligible signal. A sharp null at $\kappa = 1$ and a meaningful flip at small κ are *irreconcilable* in this model class, because κ small *means* the read subspace is low-energy. κ is therefore a coupling-null predictor with an exact null condition (read support = signal support), not a magnitude dial; the flip magnitude is governed by where the read-subspace *signal* sits relative to the reconstruction allocation, which is not a single scalar. Recorded as an honest negative (`results/G0-kappa-law.json`).

The endpoints of the ordinal prediction are fixed *analytically*, not fit:

- $\kappa \rightarrow 0$ (**maximal flip**). For direction finding the read direction is $\hat{g} = P_a^{\perp} a'(\theta_0)$, the array-manifold derivative orthogonalised against the steering vector, which is by construction orthogonal to the dominant received energy (the steering direction). For the synthetic and planted anchors the read subspace is planted orthogonal to the signal energy. Both give $\kappa \approx 0$ and the largest observed flips.
- $\kappa \rightarrow 1$ (**null**). For gradient compression under the curvature metric the read operator is the loss Hessian $H = \frac{1}{n} X^{\top} \text{diag}(p(1-p)) X$, whose high-curvature directions coincide with the high-energy directions of $X^{\top}X$; hence $\kappa \approx 1$ and, as §4 confirms, $\Delta \approx 0$ while the anti-probe stays maximal — the read operator governs the task, but the flip cannot appear.

Proposition 1 makes boundary condition (D1) an ex-ante *sign* coordinate: given the read operator and the signal covariance, κ near 1 predicts *no* flip before any code is built, which is what the gradient-compression null (§4) confirms. Remark 2.3 is equally load-bearing: the attempt to promote this to a continuous magnitude law was made and *failed* the honest test, so we do not claim it. The prospective magnitude seal we had reserved (G0-P-2026-040) is accordingly **not sealed** — registering a magnitude prediction on an unvalidated law is exactly the over-reach the protocol exists to prevent.

2.4 Recovering the read operator

One rarely has $P_{\mathcal{C}}$ in closed form. The read operator enters three ways. *Geometric* — for sensor arrays it is $\hat{g} = P_a^{\perp} a'(\theta_0)$, known analytically. *Planted* — in the synthetic and LLM-rematch anchors it is constructed from a known subspace and the prediction made blind to it. *Recovered* — on real learned representations it is obtained by the **blind probe** (Anonymous, 2026b): perturb the consumer’s input and read its output sensitivity. For a cosine-ranking consumer the margin sensitivity is $\partial \cos(a, b) / \partial b \approx \hat{a} - \cos(a, b) \hat{b}$ (the query component orthogonal to the item), and the read operator is the top- r eigenspace of $S = \langle g g^{\top} \rangle$ — low-rank for robustness, with r and the rate selected on an internal calibration split so that only a configuration that already generalises once is sealed. The blind probe is the constructive form of the identifiability face (D3): it is what makes the flip available on any real learned representation whose consumer can be queried.

3 Methods

A domain-general protocol. Every domain is tested by the same design. On the domain’s native representation we build **three codes at matched total bits per sample**: read-preserving (R), reconstruction-optimal (O, Lloyd–Max or reverse water-filling), and anti (A). The consumer is the domain’s own mature estimator, independent of the compressor. For each held-out instance we record the downstream error and the reconstruction error; the per-instance flip and anti checks are (2). All configuration is chosen on a

The flip: at matched bits the two fidelities move oppositely




Figure 1: **The flip, on two real held-out consumers.** At matched bits the read-preserving code (R) beats the reconstruction-optimal code (O) on the downstream task (left; AUROC ↑) *while* O reconstructs the raw signal better (right; relative MSE ↓). The two fidelities move in opposite directions — a code can win reconstruction and lose the task. Numbers are held-out (legal on the virgin split, §4; whale $n=2907$).

calibration split and frozen before the held-out run; each domain is pre-registered (a content hash sealed and git-timestamped before any held-out instance is scored), and outcomes are recorded regardless of sign. Matched bits are audited (allocation spread $\leq 1\text{--}2\%$); flip fractions are bootstrapped over query subsets for the retrieval and classifier consumers.

The (A2) regime diagnostic. For embedding consumers we additionally run `probe_quotient` (the compression engine’s calibration-time probe): at matched bits it applies a polar (angular), a per-channel (reconstruction-oriented), and a ZCA-whitened proxy code and reports the Spearman rank agreement of the declared consumer’s scores under each. Two summary statistics characterise the regime: `tangential_fraction` ($\approx 1 \Rightarrow$ an angular consumer) and `median_unit_displacement` ($\sqrt{2} \Rightarrow$ isotropic directions; small $\Rightarrow$ a concentrated offset cone).

Domains and data. *Arrays* — a simulated $\lambda/2$ ULA; LOCATA (Evers et al., 2020) (8cm nested ULA, OptiTrack ground truth); AV16.3 (Lathoud et al., 2004) (8-mic circular array, camera-tracked 3-D mouth); PDAR (13-element IMS seismic array, USGS catalogue back-azimuth, 34 teleseismic events; waveforms via EarthScope Consortium, 2024); RaDiCaL (Lim et al., 2021) (8-element 77 GHz FMCW virtual array). The array consumer throughout is (wideband) MUSIC (Schmidt, 1986), and the reconstruction-optimal baseline is Lloyd–Max quantization (Lloyd, 1982). *Neural* — Llama-3.2-3B (Grattafiori et al., 2024) attention keys (layers 8/16); a logistic-regression training step (loss Hessian as the read metric). *Learned representations* — CourtListener (Free Law Project, 2024) citation pairs with LaBSE (Feng et al., 2022) (cosine ranking); ETHICS commonsense (Hendrycks et al., 2021) with a bert-base classifier fine-tuned to 0.74 (frozen embeddings sit at chance); FMA music (Defferrard et al., 2017) with CLAP (Elizalde et al., 2023) PCA embeddings and a genre classifier; sperm-whale codas (Sharma et al., 2024) (DSWP) with a clan/dialect classifier on the inter-click-interval rhythm. Held-out sets are disjoint recordings/events/opinions/ codas wherever obtainable.

Table 1: The twelve-domain sweep. Tier $\in \{\text{sealed confirmation, partial, sealed null, (A2) verdict}\}$. Every row resolves to a sealed pre-registration and a result artifact (Appendix A, Appendix B).

Domain	Consumer / read operator	Held-out outcome	Tier
Synthetic ULA	beamformer-ML + root-MUSIC / $\hat{g}$	flip 6/6 (two estimators)	sealed conf.
LLM keys (Llama-3.2-3B)	softmax attention / blind probe	worse-arm 16/16, recon tied	sealed conf.
LOCATA acoustic	wideband MUSIC / array manifold	flip 11/13, recon-trade 13/13	sealed conf.
AV16.3 acoustic	wideband MUSIC / array manifold	flip 74%, recon-trade 100%	sealed conf.
PDAR seismic	wideband MUSIC baz / array manifold	flip 76%, recon-trade 100%	sealed conf.
Legal retrieval	cosine ranking / blind probe	R 0.796 > O 0.780 (virgin split)	sealed conf.
Whale codas	clan/dialect classifier / blind probe	R 0.934 > O 0.883, flip 300/300	sealed conf.
RaDiCaL radar	MUSIC DOA / virtual array	flip 25/25, anti 60% (data-limited)	partial
Gradient optim.	curvature metric / loss Hessian	anti 300/300, flip 27%	sealed null
Music genre	genre classifier / weight	polar 0.994 > per-ch 0.956	(A2) verdict
Moral judgment	fine-tuned classifier / weight	polar preferred (once trained)	(A2) verdict




Figure 2: **The flip across nine sealed domains and six physics.** Held-out flip fraction (fraction of bootstrap query-subsets / held-out instances on which R beats O) with Wilson 95% intervals, coloured by tier. Seven sealed confirmations span synthetic, LLM-attention, acoustic, seismic, retrieval, and bioacoustic consumers; radar is a data-limited partial; gradient compression is the sealed coupling null (D1), sitting at chance. Counts k/n at right; the LLM-keys row is a reconstruction-invisible worse-arm discrimination (16/16 heads).

4 Results

Table 1 is the spine; each row wears its tier. Figure 1 shows the dissociation on two consumers; Figure 2 places it across the sweep.




Figure 3: **Recovering the read operator makes the flip transfer.** On legal-citation retrieval an *estimated* read operator overfits and reconstruction-optimal (O) beats read-preserving (R) on held-out data (G0-P-2026-035, shaded miss). Replacing it with the GO-1 blind probe — the only change — reverses the verdict on the identical held-out set (G0-P-2026-036) and again on a fully virgin, opinion-disjoint split (G0-P-2026-039). The 035 miss was a read-operator-identifiability failure (D3), not a failure of the flip.

The flip reproduces across three physics. On the physical arrays it holds under the mature subspace estimator and against independent ground truth: LOCATA **11/13** held-out recordings with the reconstruction-trade at **13/13** (Lloyd reconstructs better every time yet is downstream-worse); AV16.3 **74%** of held-out windows with recon-trade **100%** and anti-worst **76%**, against camera-tracked angle on *disjoint* recordings; PDAR seismic back-azimuth **76%** of held-out events, recon-trade **100%**, anti **76%**, against catalogue ground truth. The synthetic anchor flips 6/6 under two independent estimators; the Llama-3.2-3B key consumer shows the read-operator-projected error predicting the softmax-KL loser on **16/16** heads while exactly-tied reconstruction cannot (a reconstruction-invisible *worse-arm* discrimination, distinct from a two-consumer inversion; §6). The battery’s pre-registered prediction — a flip in ≥ 2 of its three core prospective domains — is met by the two confirmed arrays.

Recovering the read operator makes the flip transfer, twice. On real legal-citation retrieval an *estimated* read operator (document covariance) overfits: on held-out opinions reconstruction-optimal *beats* read-preserving (AUROC 0.773 vs 0.757), a registered miss (G0-P-2026-035, degeneracy D3). Replacing it with the blind-probe read operator — the only change — reverses the verdict on the identical held-out set (0.779 vs 0.771; flip 200/200 bootstraps) while reconstruction-optimal still reconstructs $2.6\times$ better (G0-P-2026-036). Because that held-out split had itself produced the 035 verdict and the margin was thin, we re-ran the *frozen* recipe on a fully virgin, opinion-disjoint split (evaluation opinions never previously scored), sealed before any virgin instance was embedded (G0-P-2026-039): all four bars pass and more cleanly — **R 0.796** > **O 0.780**, a margin double that of 036, with R edging even the uncompressed embedding (0.794, nuisance-stripping denoises the ranking), recon $O\ 0.223 \leq R\ 0.561$, flip 200/200, anti 200/200. Moral judgment shows the consumer precondition (D2): frozen embeddings classify at chance (0.58–0.60 vs 0.63 majority) and the flip is ill-posed; a classifier fine-tuned to 0.74 makes the angular code preserve the class logit (0.994 vs 0.956 at 1 bit).

The whale flip is sealed, not exploratory. Sperm-whale coda dialect is the biologically most distant domain and the load-bearing cell of the regime map (isotropic directions, yet correlated channels). It was promoted from an exploratory (A2) verdict to a sealed flip (G0-P-2026-038): with the clan classifier as consumer and the blind margin probe as read operator, on held-out disjoint codas ($n = 2907$) read-preserving

AUROC **0.934** > **0.883** reconstruction-optimal, flip **300/300**, anti **300/300**, while reconstruction-optimal reconstructs the click-interval features $2\times$ better. The map’s novel axis now rests on a confirmation.

Two boundary conditions bound the claim. Gradient compression under a curvature metric is the *coupling* boundary (D1, $\kappa \approx 1$): the anti-probe is worst on **300/300** held-out steps — the Hessian read operator genuinely governs the update — yet the flip sits at chance (**27%**), because the high-curvature and high-gradient-energy directions coincide, so reconstruction already protects the read directions. No read-operator recovery can rescue it; the operator was already exact. Radar is the data limit: the flip and reconstruction-trade hold **25/25** on real 77 GHz FMCW returns, but the anti control reaches only 60% on the 8-element array and the disjoint-recording held-out is inaccessible (the full set is 310 GB) — recorded as a partial, not inflated to a confirmation.

5 The regime map: two axes, not one

Characterising every embedding consumer with `probe_quotient` places them on two *independent* axes — direction concentration and cross-channel correlation — and the winning code is set by both (Table 2). Music (isotropic, low correlation) is preserved by the generic angular code and reconciles a pre-existing 6σ music sign-flip; legal (concentrated, channel-structured) needs the per-channel or blind-probe code; post-RoPE keys and sperm-whale codas both need the whitened code, but for different reasons — the keys are concentrated-correlated, the codas isotropic-correlated (their inter-click intervals are temporally correlated). Whale is the decisive case: isotropic like music yet whitening-dependent like the keys, so a single concentration statistic cannot predict it. Indeed a pre-registered `unit_displacement` threshold is *refuted* on its first out-of-sample test (G0-P-2026-037, ledger NEG-14): synthetic concentrated embeddings sit below the threshold yet the generic polar code wins. The reliable diagnostic is the full end-to-end probe run per domain at low bits.

Table 2: The two-axis regime map. Cells give the winning code; the empty cell (concentrated / low-correlation) is argued structurally empty in Appendix A.

	isotropic (<code>unit_disp</code> $\approx \sqrt{2}$)	concentrated (<code>unit_disp</code> small)
low correlation	music, moral $\rightarrow$ polar	(<i>structurally rare</i>)
cross-channel correlated	whale $\rightarrow$ whitened	keys $\rightarrow$ whitened
channel-structured	—	legal $\rightarrow$ per-channel / blind-probe

6 Discussion

What is established, and at what tier. The load-bearing result is **seven sealed confirmations** — synthetic, LLM attention, LOCATA, AV16.3, PDAR, legal (double-confirmed on a virgin split), and whale — each scored on held-out data against independent ground truth under a pre-registered protocol. Around them the sweep carries, and labels, weaker evidence: one **partial** (radar, data-limited), one **sealed null** (gradient compression), and two (**A2**) **verdicts** (music, attention keys). Each wears its tier. We draw this line deliberately, because it is what lets a reader trust the confirmations without re-running them: the same registration machinery that seals a confirmation also makes an exploratory row legible as exploratory. The full registry accounting — every registration identifier, its disposition, and the seal-before-run commit ordering — is Appendix B.

The flip and the (A2) verdict are not the same claim. We reserve *flip* for the two-metric dissociation (2) established by the sealed protocol. The (**A2**) **verdict** is a distinct, one-metric quantity: which quantizer family best preserves the consumer’s rank ordering at matched bits. A verdict is *evidence consistent with* the flip but does not by itself establish the reconstruction-trade. Two rows report verdicts; seven report flips; the terms are kept apart throughout. (A separate one-word disambiguation applies to the LLM-keys

anchor: it is a reconstruction-invisible *worse-arm* discrimination on a single consumer, not a two-consumer verdict *inversion*.)

The failure taxonomy is what makes this a theory. The four degeneracies of P_C (§2.2) bound the claim where it fails. Identifiability is fixable by recovery; coupling is a wall; precondition needs a working consumer; mechanism-absence is a refutation. The gradient-compression null is coupling’s empirical face, and the κ law (§2.3) turns that wall into a graded, ex-ante prediction rather than a category. The refuted single-statistic shortcut is the same discipline applied to the regime map: a predictor pre-registered and killed on its first out-of-sample test, which is what forces the regime to be characterised, not asserted.

The information-theoretic frame, and convergent evidence. That reconstruction fidelity is the special case $P_C = I$ places the flip inside rate–distortion theory (Shannon, 1959; Cover & Thomas, 2006) as its consumer-relative generalisation (Anonymous, 2026a). Independent arrival at the same object is telling: a bioacoustics decoder-robustness index — the public name for the read- operator-distortion measure whose provenance we footnote¹ — measures precisely read-operator distortion (invariant acoustic transforms leave the decoder’s read unchanged; stress transforms move it), i.e. the same methodology reached from a different field.

Limitations. The two reconciliations are (A2) verdicts, not sealed flips; radar is data-limited; each domain uses a single model and dataset; the audio-level decoder-robustness measurement was not run. And the κ coordinate delivers less than we first hoped: it robustly predicts the coupling *null* ($\kappa \rightarrow 1 \Rightarrow$ no flip) but not the flip magnitude (Remark 2.3), a limitation we found by testing the stronger claim and reporting its failure rather than the claim.

Outlook. The operational reading is direct: for compression under a known consumer, *measure the regime and adapt the code* — certify what the consumer reads rather than promise reconstruction. There is no universal code; there is a diagnostic, on two axes, and a selector that reads it — and, for the boundary, a single alignment number that says in advance how much the flip will buy.

References

- Anonymous. A consumer-relative rate–distortion theory (companion submission), 2026a. Under review; author omitted for double-blind review.
- Anonymous. Identifiability of a consumer’s read operator from black-box queries (companion submission), 2026b. Under review; author omitted for double-blind review.
- Thomas M. Cover and Joy A. Thomas. *Elements of Information Theory*. Wiley, 2nd edition, 2006.
- Michaël Defferrard, Kirell Benzi, Pierre Vandergheynst, and Xavier Bresson. FMA: A dataset for music analysis. In *18th International Society for Music Information Retrieval Conference (ISMIR)*, 2017.
- EarthScope Consortium. IRIS/EarthScope data management center: PDAR (ims array) waveform data via FDSN. <https://www.earthscope.org>, 2024.
- Benjamin Elizalde, Soham Deshmukh, Mahmoud Al Ismail, and Huaming Wang. CLAP: Learning audio concepts from natural language supervision. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2023.
- Christine Evers, Heinrich W. Löllmann, Heinrich Mellmann, Alexander Schmidt, Hendrik Barfuss, Patrick A. Naylor, and Walter Kellermann. The LOCATA challenge: Acoustic source localization and tracking. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 28:1620–1643, 2020. doi: 10.1109/TASLP.2020.2990485.

¹Introduced under a different name in an earlier ethics setting and renamed the Decoder Robustness Index (DRI) as the public index name in the bioacoustic setting; the original name (which identifies the authors) is restored in the non-anonymized version.

- Fangxiaoyu Feng, Yinfei Yang, Daniel Cer, Naveen Arivazhagan, and Wei Wang. Language-agnostic BERT sentence embedding. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL)*, pp. 878–891, 2022. doi: 10.18653/v1/2022.acl-long.62.
- Free Law Project. CourtListener. <https://www.courtlistener.com>, 2024. Legal opinion corpus and citation graph.
- Aaron Grattafiori, Abhimanyu Dubey, et al. The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.
- Dan Hendrycks, Collin Burns, Steven Basart, Andrew Critch, Jerry Li, Dawn Song, and Jacob Steinhardt. Aligning AI with shared human values. In *International Conference on Learning Representations (ICLR)*, 2021.
- Guillaume Lathoud, Jean-Marc Odobez, and Daniel Gatica-Perez. AV16.3: An audio-visual corpus for speaker localization and tracking. In *Machine Learning for Multimodal Interaction (MLMI 2004)*, LNCS 3361, pp. 182–195. Springer, 2004.
- Teck-Yian Lim, Spencer A. Markowitz, and Minh N. Do. RaDICAL: A synchronized FMCW radar, depth, IMU and RGB camera data dataset with low-level FMCW radar signals. *IEEE Journal of Selected Topics in Signal Processing*, 15(4):941–953, 2021. doi: 10.1109/JSTSP.2021.3061270.
- Stuart P. Lloyd. Least squares quantization in PCM. *IEEE Transactions on Information Theory*, 28(2): 129–137, 1982.
- Ralph O. Schmidt. Multiple emitter location and signal parameter estimation. *IEEE Transactions on Antennas and Propagation*, 34(3):276–280, 1986.
- Claude E. Shannon. Coding theorems for a discrete source with a fidelity criterion. In *IRE National Convention Record, Part 4*, pp. 142–163, 1959.
- Pratyusha Sharma, Shane Gero, Roger Payne, David F. Gruber, Daniela Rus, Antonio Torralba, and Jacob Andreas. Contextual and combinatorial structure in sperm whale vocalisations. *Nature Communications*, 15:3617, 2024. doi: 10.1038/s41467-024-47221-8.

A Per-domain detail

Each row of Table 1 expands to a per-domain entry recording the arms, the sample size n , the sealed bars and the realised statistics with their binomial nulls, the read-operator provenance (geometric / planted / recovered), and the seal and result commit hashes. These reproduce the ledger’s Gate-B rows and the statistical audit (every reported percentage with its n , its bar, and its null); they are omitted here for length and provided in the supplementary registry. The empty regime cell (concentrated / low-correlation) is argued structurally rare: a representation whose directions are concentrated into a shared offset cone is, generically, one whose channels have acquired cross-channel correlation through that shared component — so concentration without correlation is the off-diagonal that the data do not populate.

B Registry accounting

The “no file drawer” claim rests on the registration sequence having no unexplained holes. The supplementary registry accounting gives every G0-P-2026 identifier (001–039) a disposition — reported in this sweep, reported in a companion paper (the COST theorem harnesses), superseded by a named successor (each miss still reported), or never assigned (the two numbering skips, verified absent from the full git history) — together with the commit-ordering audit showing each sealed row’s registration is a strict git ancestor of its result. The supersession chains (035 → 036 → 039, 032 → 033, 020 → 021, and the earlier core iterations) carry every miss at equal prominence with its correction.

C The κ measurement (results/G0-kappa-law.json)

We computed κ (Definition 1) for the two analytic anchors ($\kappa \approx 0.006$ for the synthetic-ULA read $\hat{g} = P_a^\perp a'$; $\kappa \approx 0.80$ for the gradient Hessian on `load_digits`), and ran a controlled synthetic sweep dialing κ from 0 to 1 with real matched-bit quantization to test whether the flip magnitude Δ is monotone in κ . It is not: Δ is peaked rather than monotone (no tested normalization — Δ , $\Delta/\text{tr}(\bar{P}_{\mathcal{C}}\Sigma_x)$, or $d_{\mathcal{C}}(\text{O})/d_{\mathcal{C}}(\text{R})$ — is monotone), because a low-energy read subspace carries little signal and because the result depends on the quantizer-range regime (Remark 2.3). The *ordinal* content survives — the $\kappa \approx 0.80$ gradient anchor is the empirical null (flip 27%), the $\kappa \approx 0.006$ array read is a large flip. The faithful coupling salvage (Remark 2.3) then pins the null to an exact condition (read support = signal support) and shows why no magnitude dial exists in this model class. The prospective magnitude seal **G0-P-2026-040** was **not** taken: the law it would have tested did not survive validation. A magnitude law matched to the quantizer regime remains open.